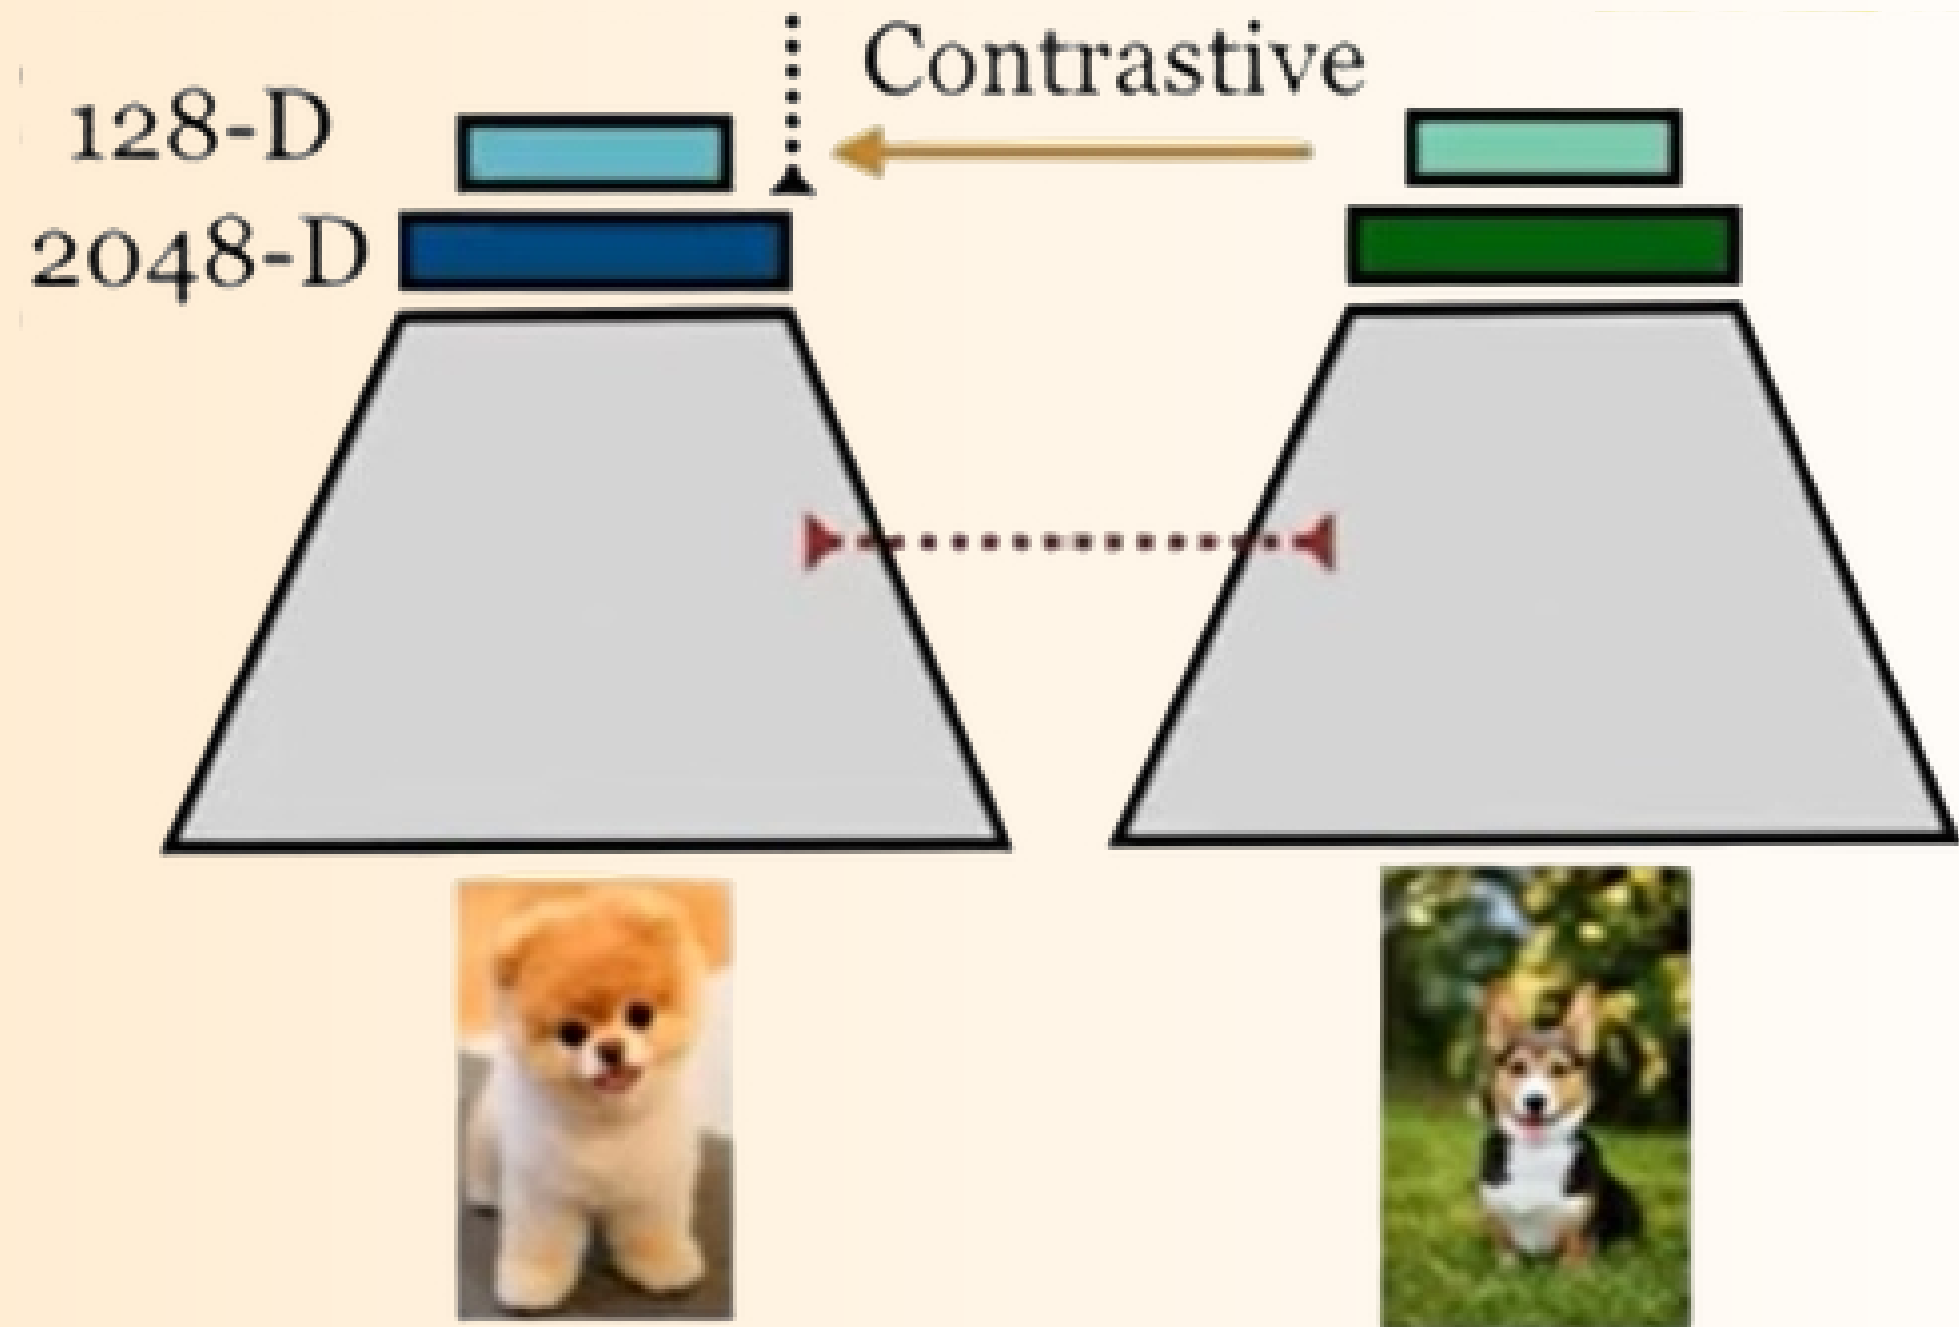
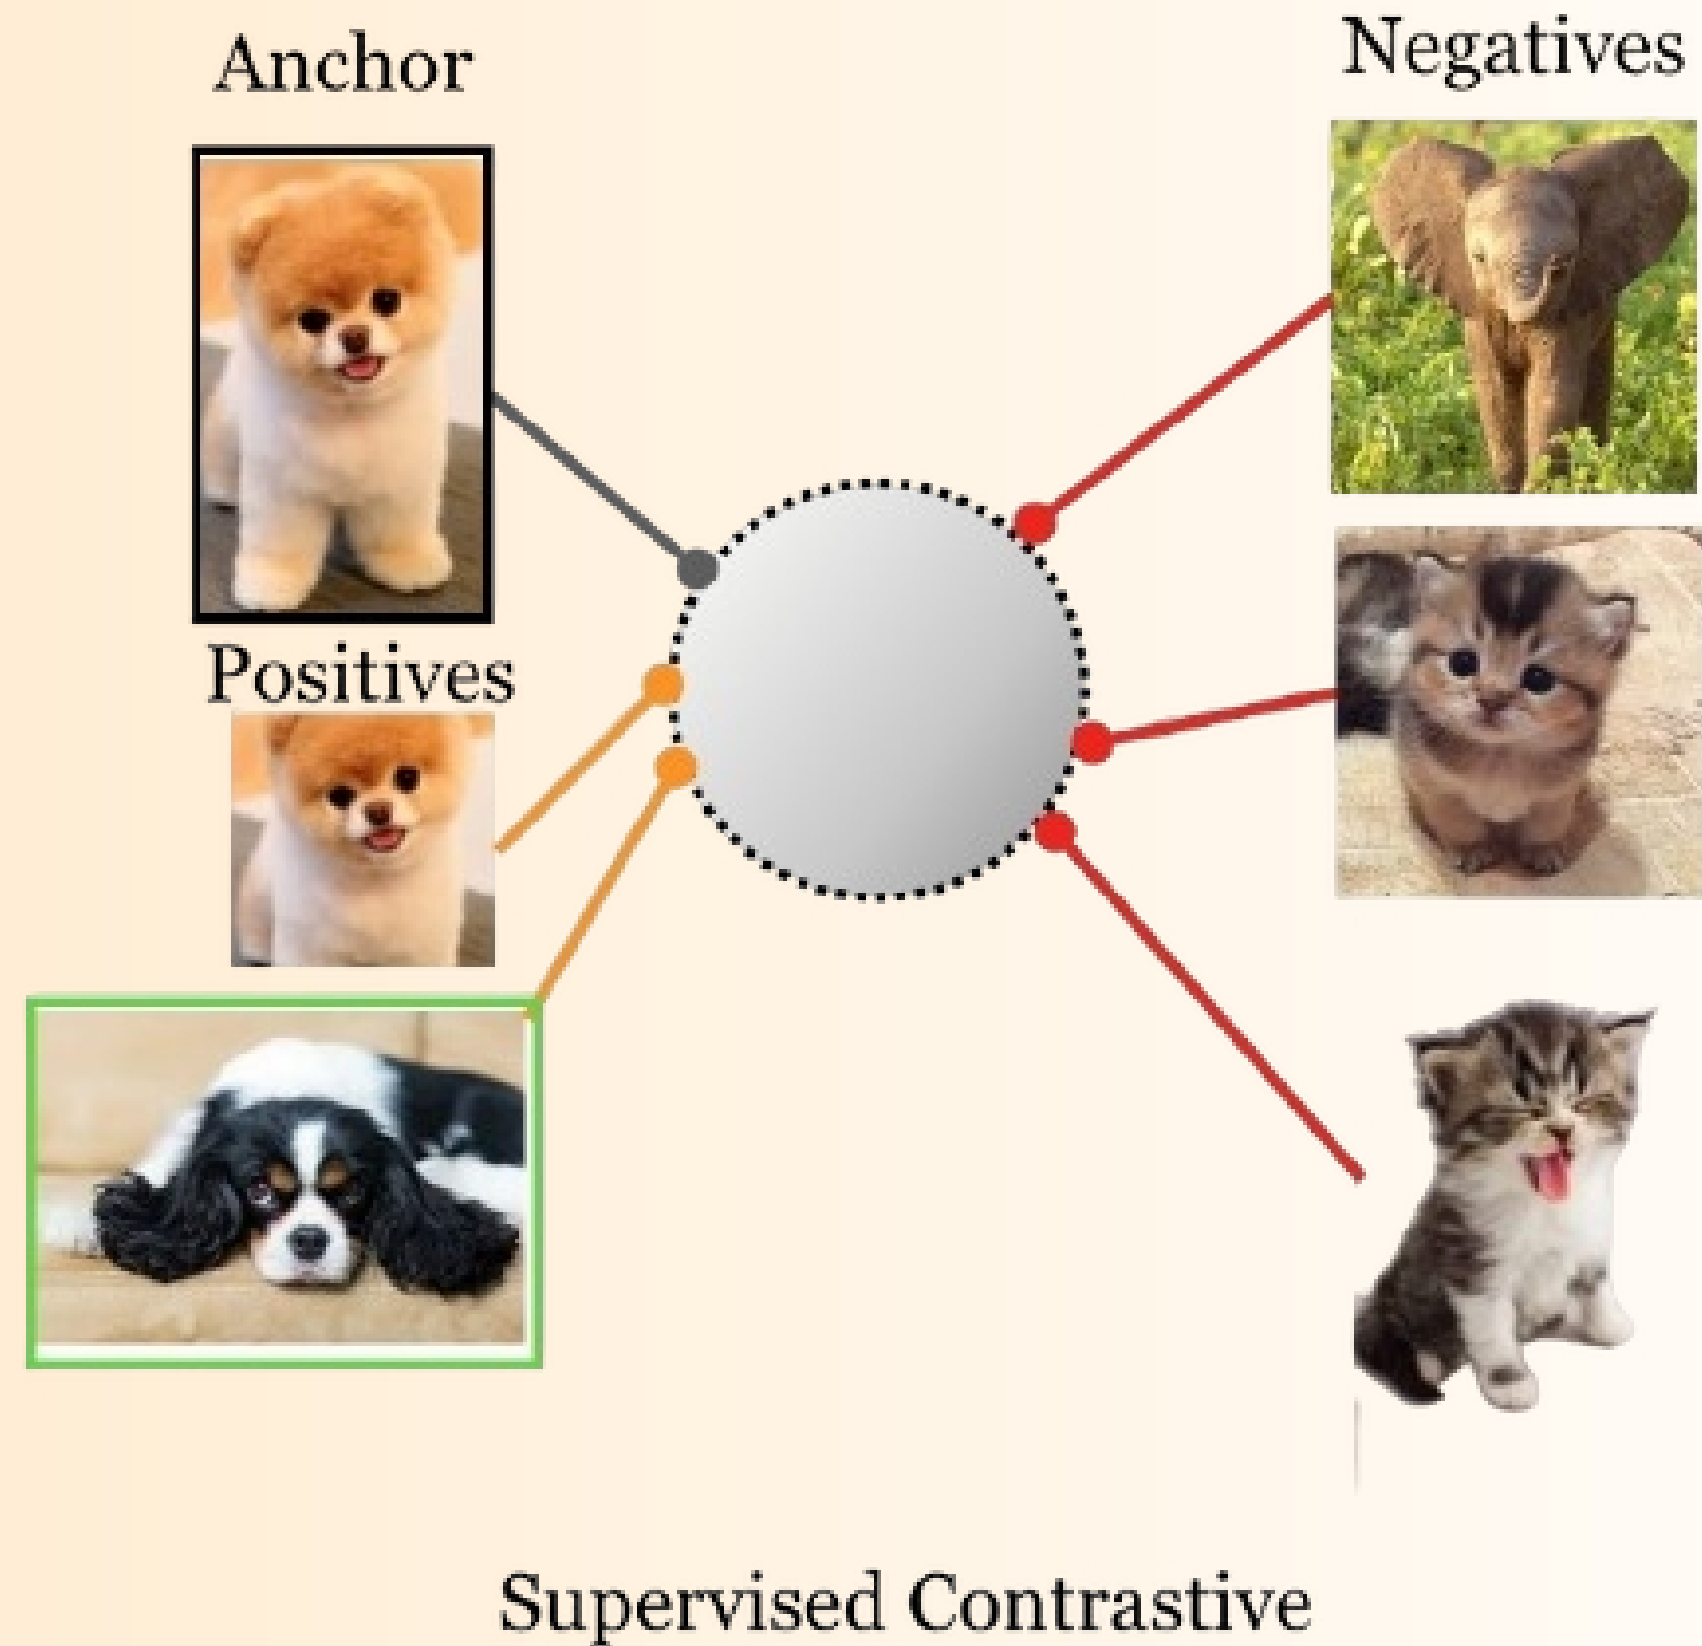


TRACK - 1

Supervised Contrastive Learning




Why Supervised Contrastive Learning?

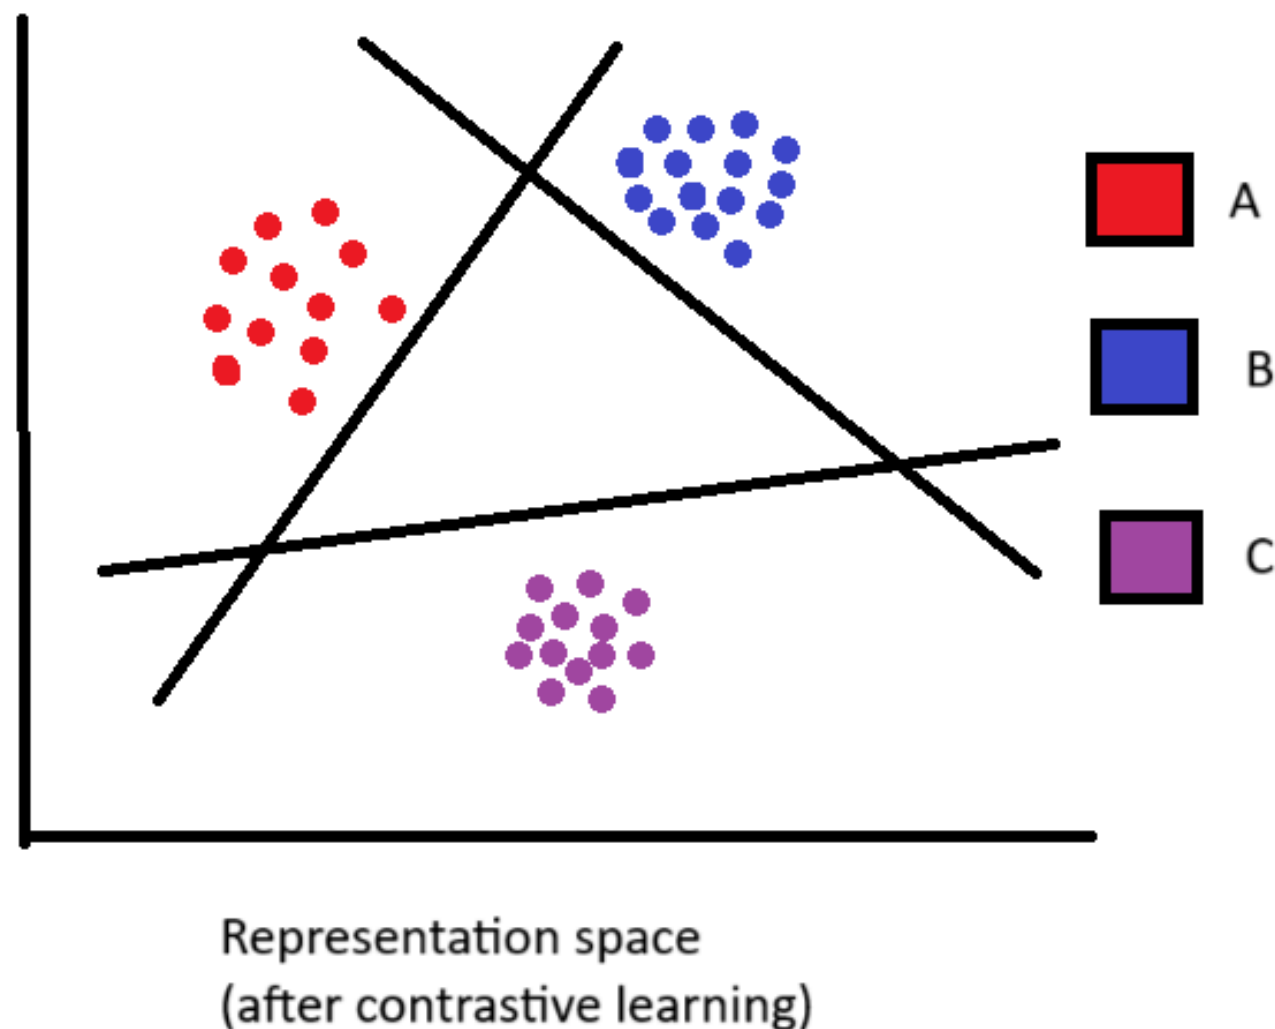


Problems Without Contrastive Training

- No clear demarcation between classes
- Difficult to be classified by linear classifiers



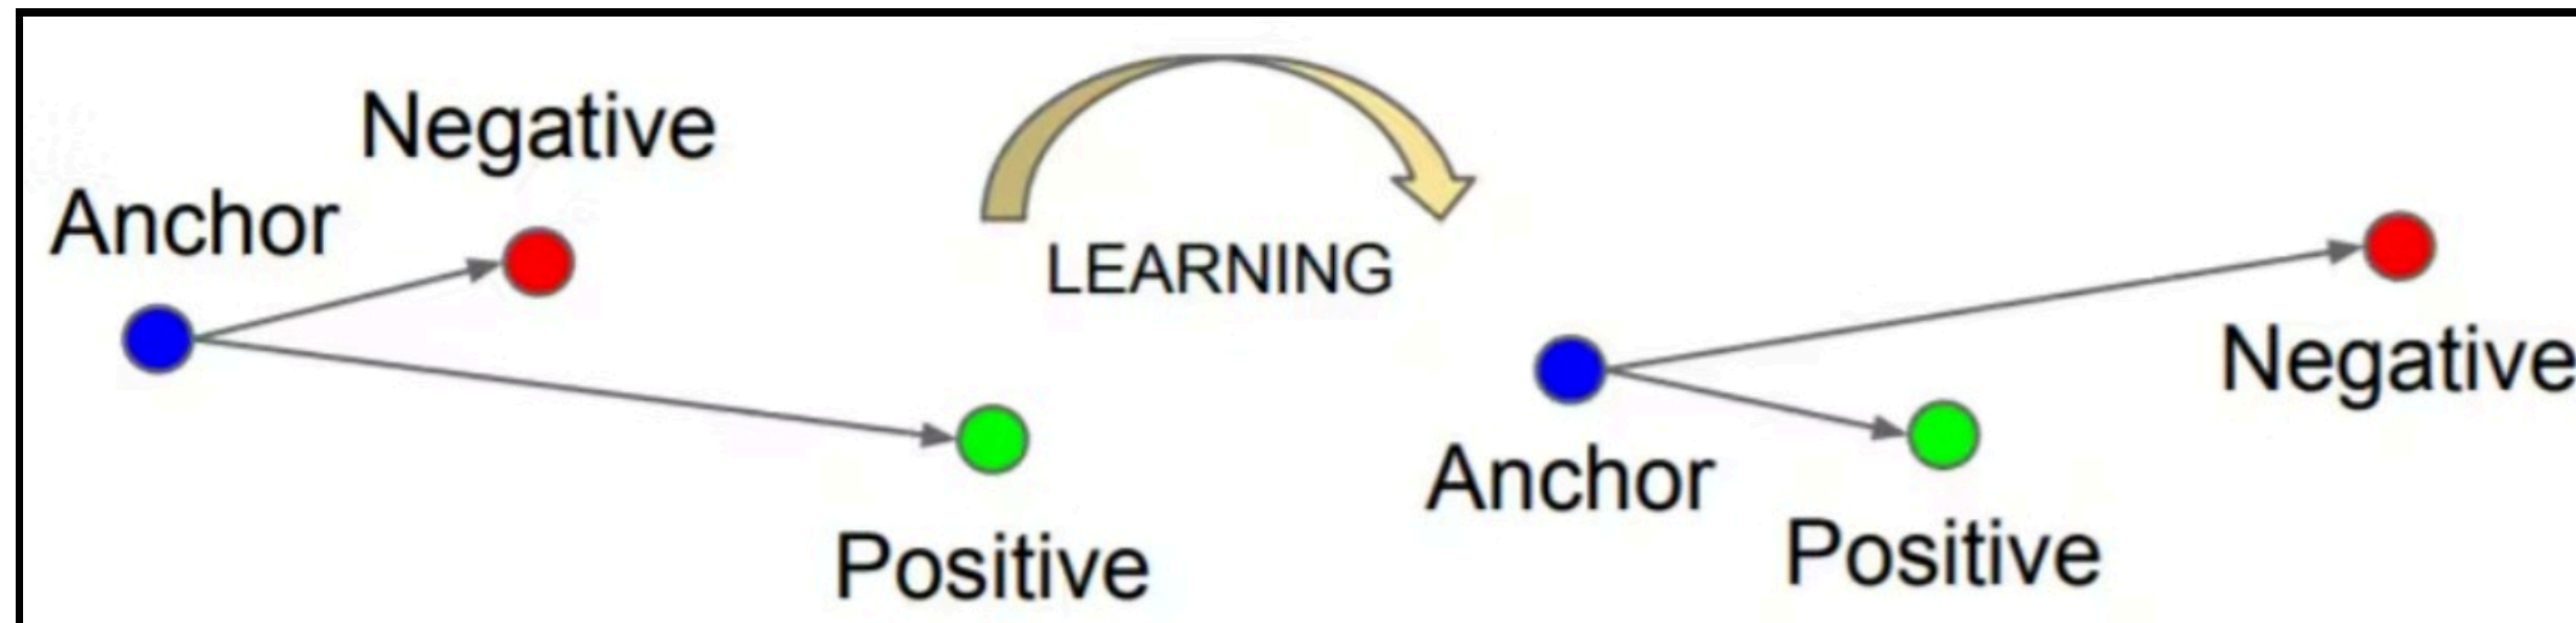

Representation space after contrastive pre-training



- After contrastive separation, the data gets fragmented into clusters in the representation space.
- This being clustered and separated makes it easier for linear classifiers.
- Thus, it gives better accuracy.
- It also requires less number of epochs for training.

How does it do that?

- We pick an anchor and push all positive samples towards it.
- We simultaneously “push” all the negative samples away from our anchor in the representation space.
- This clusters similar data together.



Contrastive Loss Function:

For each anchor, the loss tries to:

- Pull its embedding z_i closer to all of its positive samples Z_p (same class).
- Push z_i away from all other samples Z_a in the batch, regardless of class.

This is done via a log-softmax term, promoting high similarity for positive pairs and low similarity for others.

$$\mathcal{L}_{in}^{sup} = \sum_{i \in I} \mathcal{L}_{in,i}^{sup} = \sum_{i \in I} -\log \left\{ \frac{1}{|P(i)|} \sum_{p \in P(i)} \frac{\exp(z_i \cdot z_p / \tau)}{\sum_{a \in A(i)} \exp(z_i \cdot z_a / \tau)} \right\}$$

I : Set of indices for the anchor samples in the batch.

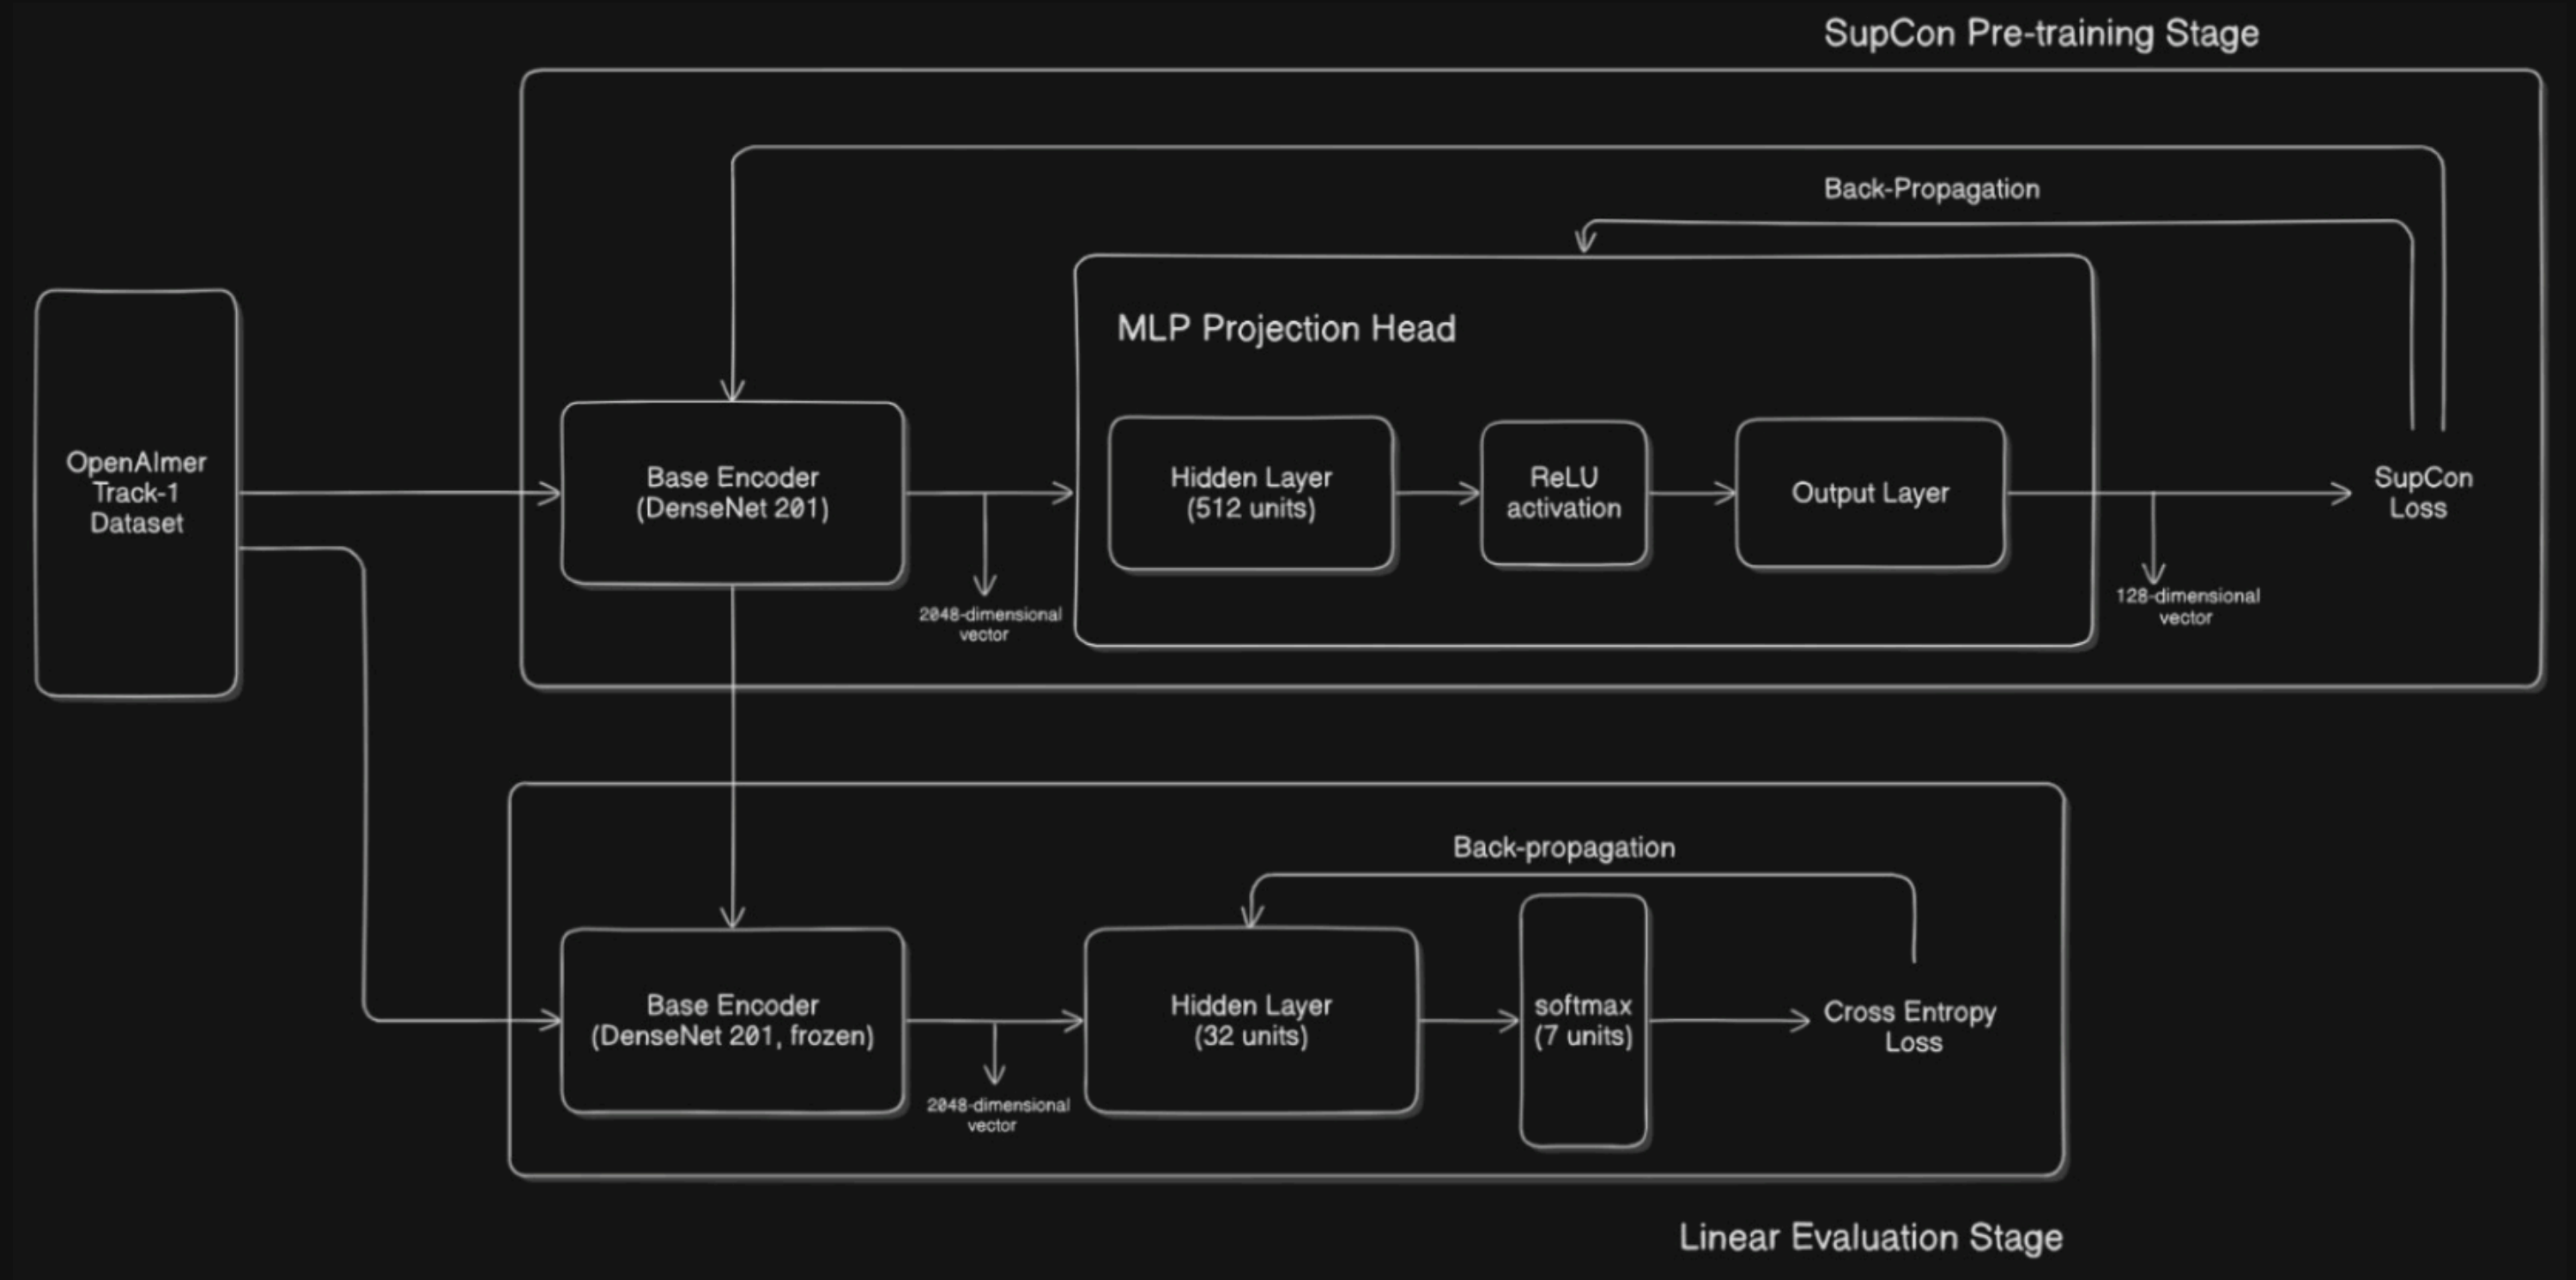
z_i : Normalized embedding of the anchor sample i .

$P(i)$: Set of indices of positives for anchor i (same class).

$A(i) = I \setminus \{i\}$: Set of all other samples except i (used for the denominator).

τ : Temperature parameter (scales similarity) — usually $\tau \in (0, 1]$.

PIPELINE




Architecture Overview



- **Two-stage Framework:** i) SupCon Pre-training Stage ii) Linear Evaluation Stage.
- **Base Encoder: DenseNet201 Architecture**
Used in both stages (trainable during SupCon stage, frozen during evaluation stage).
- **SupCon Pre-training stage:**
The DenseNet201 encoder extracts features from the input images.
Features passed through MLP Projection Head.
Encoder trained using Supervised Contrastive Loss.
- **Linear Evaluation Stage:**
Pre-trained encoder is frozen, its outputs are used as features
Classifier (softmax activation) is trained using Cross Entropy Loss for final classification
- **Back-propagation Strategy:**
In pre-training: full back-propagation through encoder + projection head.
In evaluation stage: back-propagation only through linear classifier.

DenseNet201 Model Statistics:



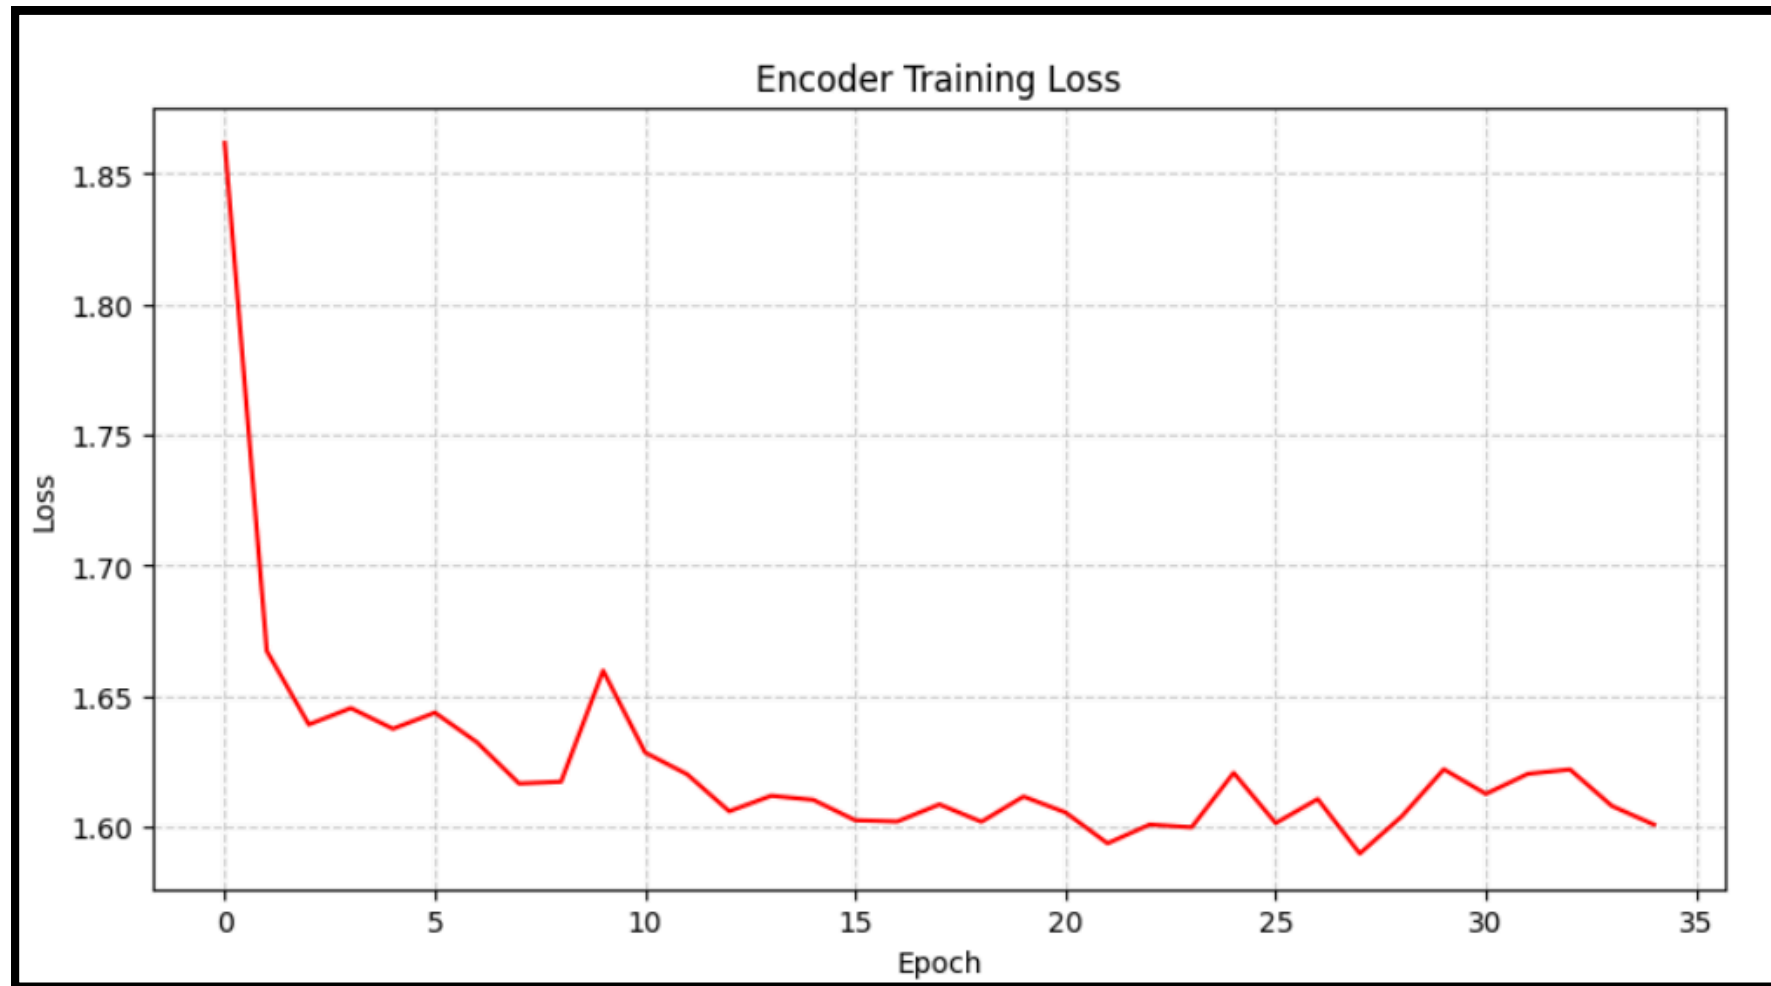
- F1-Score: 0.9948
- Accuracy: 0.9936
- Precision: 0.9944
- Recall: 0.9952



t-SNE Plots



Encoder and Classifier Performance




Other Models Used:

- 1) ResNet50
- 2) InceptionV3
- 3) InceptionResNetV2
- 4) DenseNet169

Ablation Table (CCE)					
Sl. No	Architecture	Performance Metrics			
		F1	Accuracy	Precision	Recall
1	ResNet50	0.966	0.9617	0.9637	0.9711
2	DenseNet201	0.9819	0.9836	0.9811	0.983
3	InceptionV3	0.954	0.97	0.9614	0.9479
4	InceptionResNetV2	0.9763	0.9807	0.9789	0.9738
5	DenseNet169	0.9858	0.9864	0.9877	0.9841

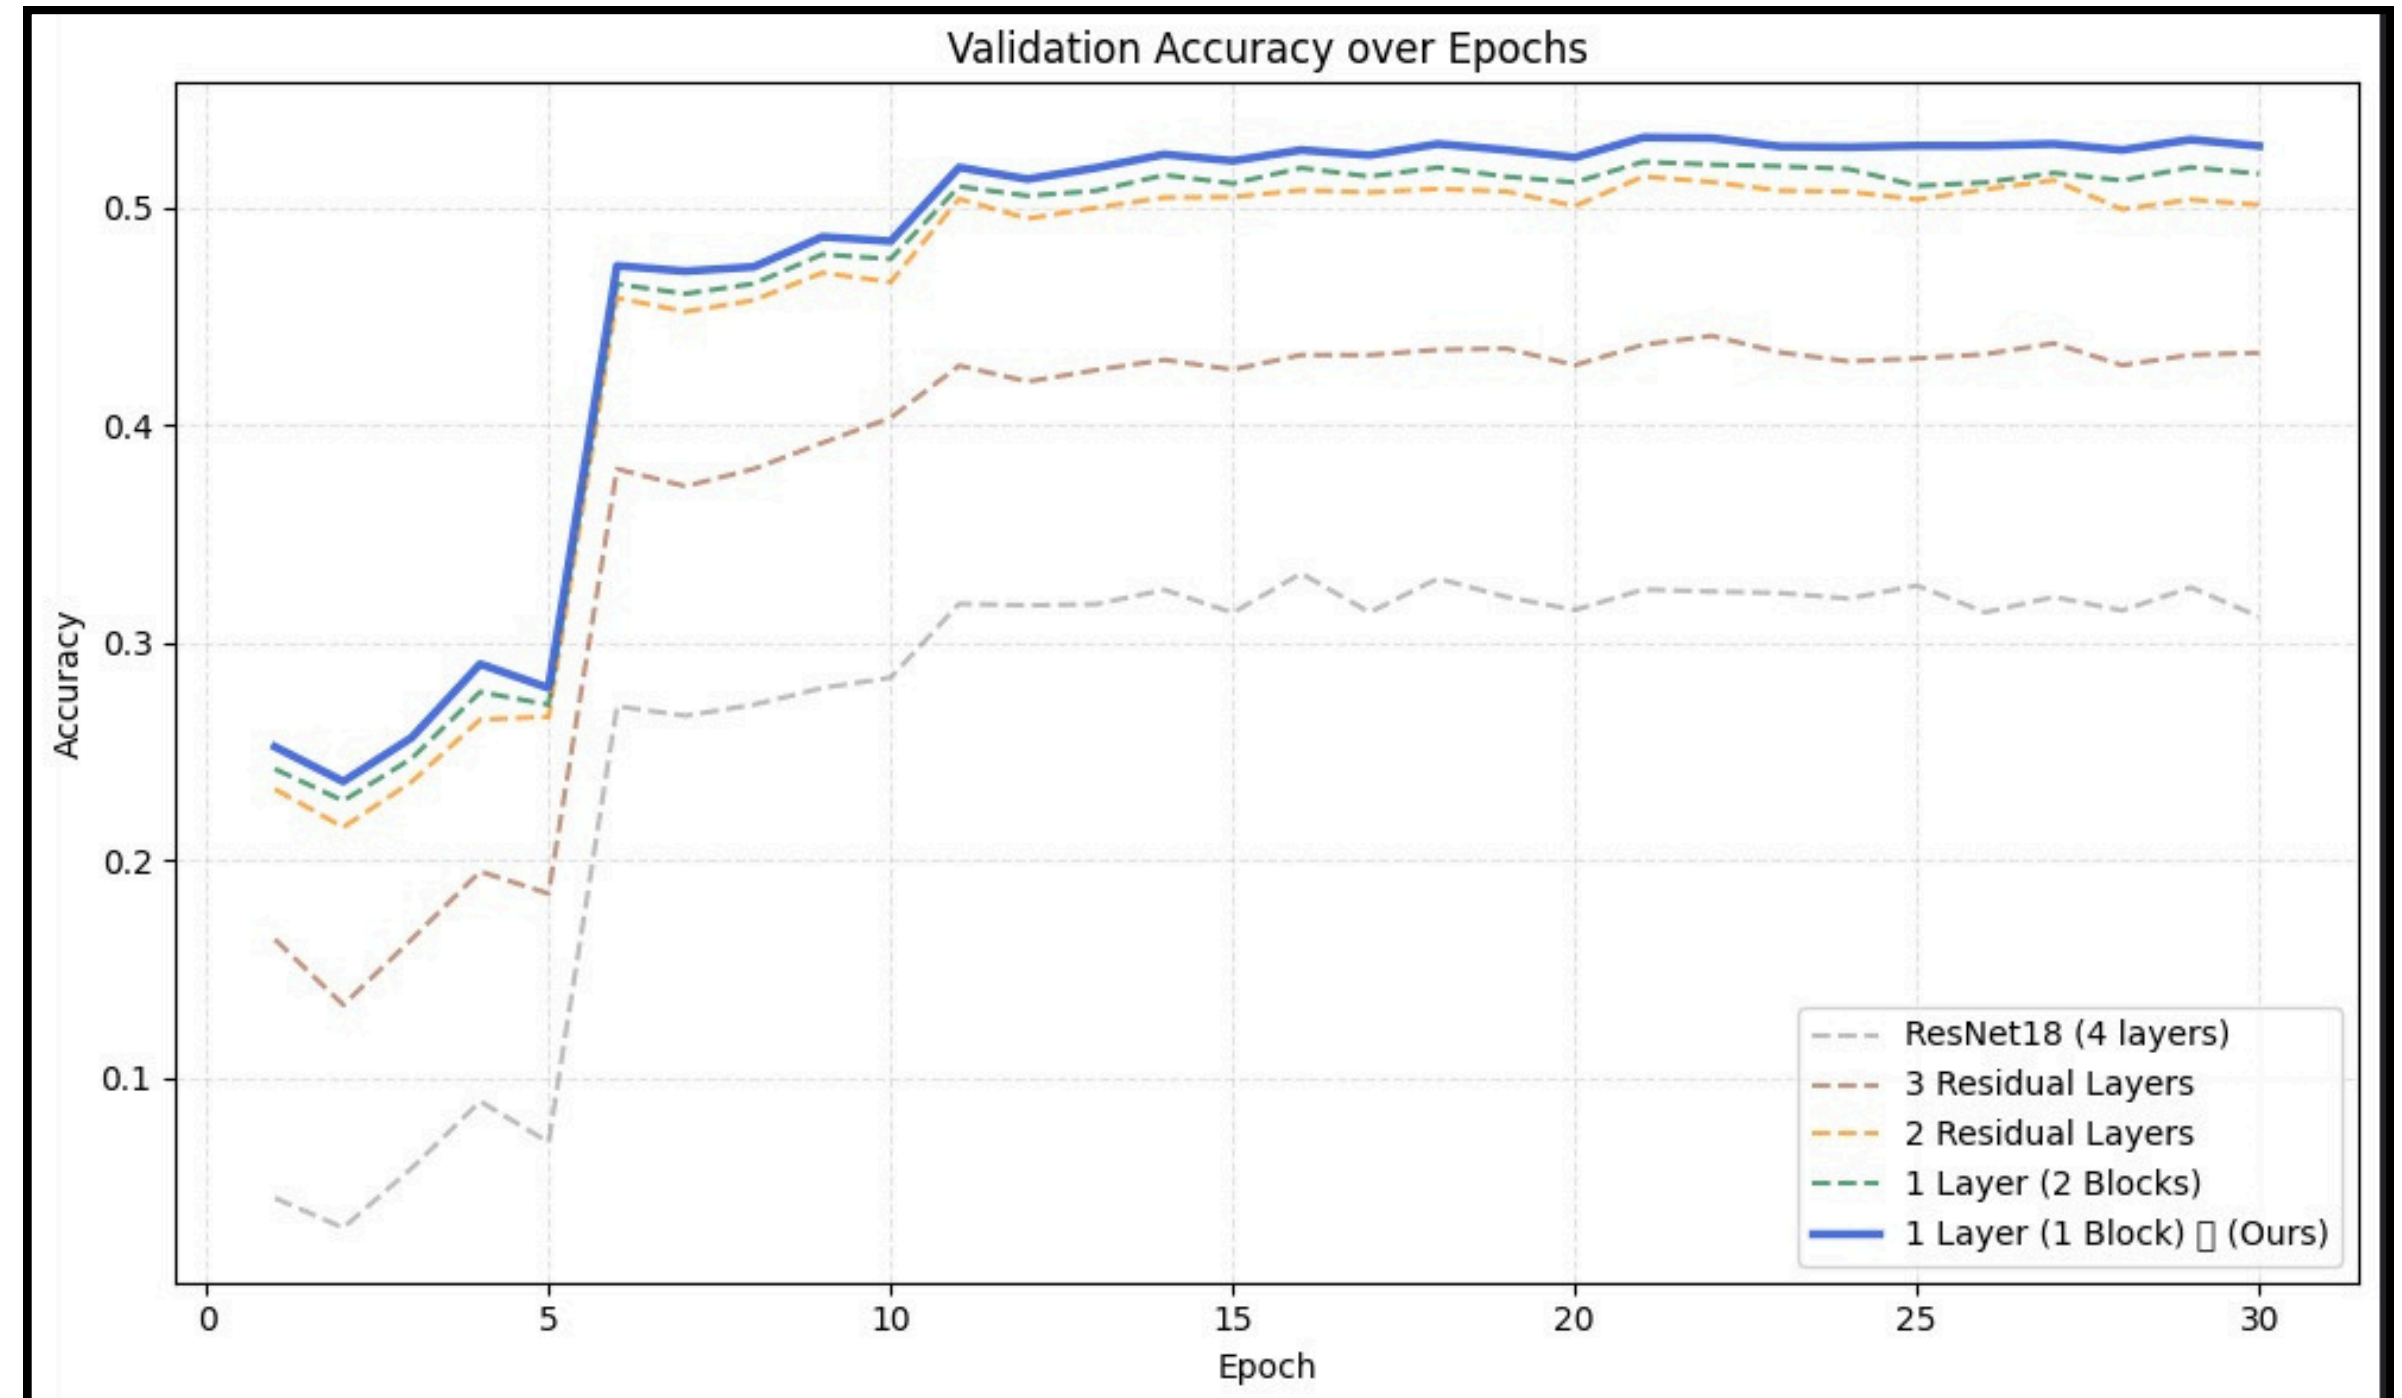
Ablation Table (SupCon)					
Sl. No	Architecture	Performance Metrics			
		F1	Accuracy	Precision	Recall
1	ResNet50	0.9856	0.9617	0.9637	0.9711
2	DenseNet201	0.9948	0.9936	0.9944	0.9952
3	InceptionV3	0.9732	0.9857	0.974	0.9724
4	InceptionResNetV2	0.9824	0.9833	0.9889	0.9838
5	DenseNet169	0.9902	0.9891	0.9931	0.9917

L1 Unstructured Weighted Pruning




Lottery Ticket Hypothesis Style Pruning - Post Training Pruning


Comparison of Final Results



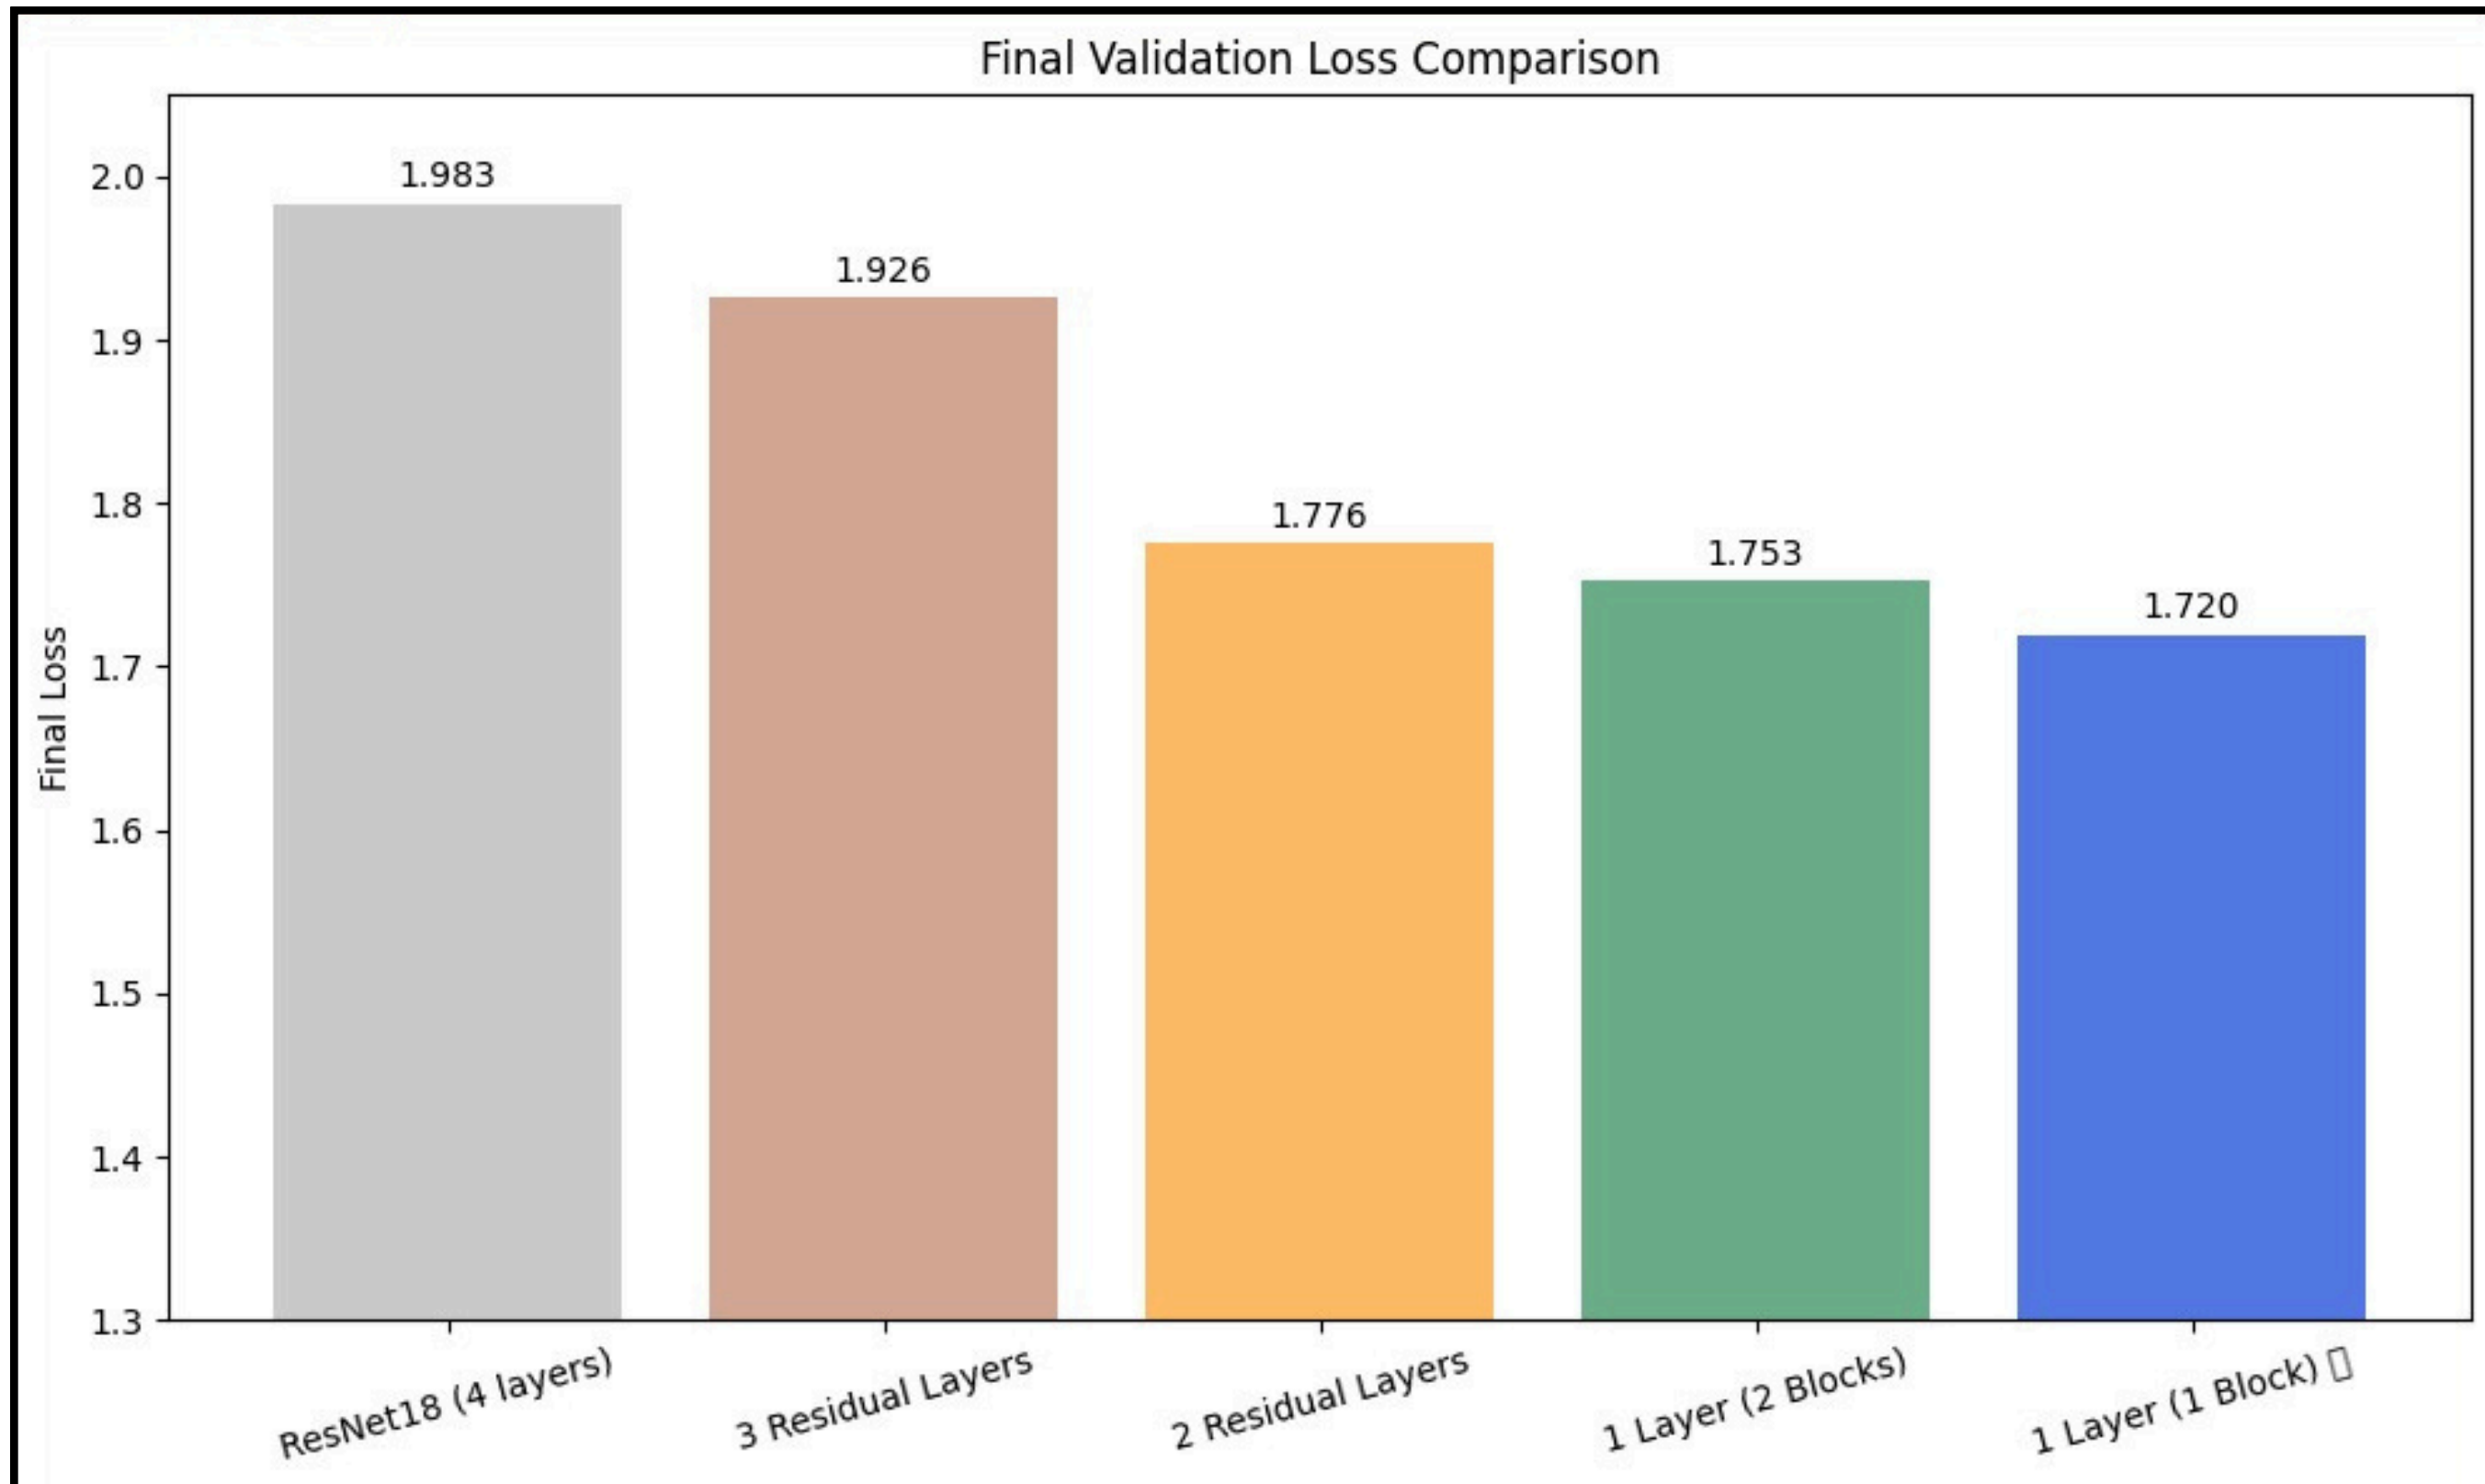
Comparison of Final Results



Final Validation Accuracy Comparison



Final Validation Accuracy Comparison



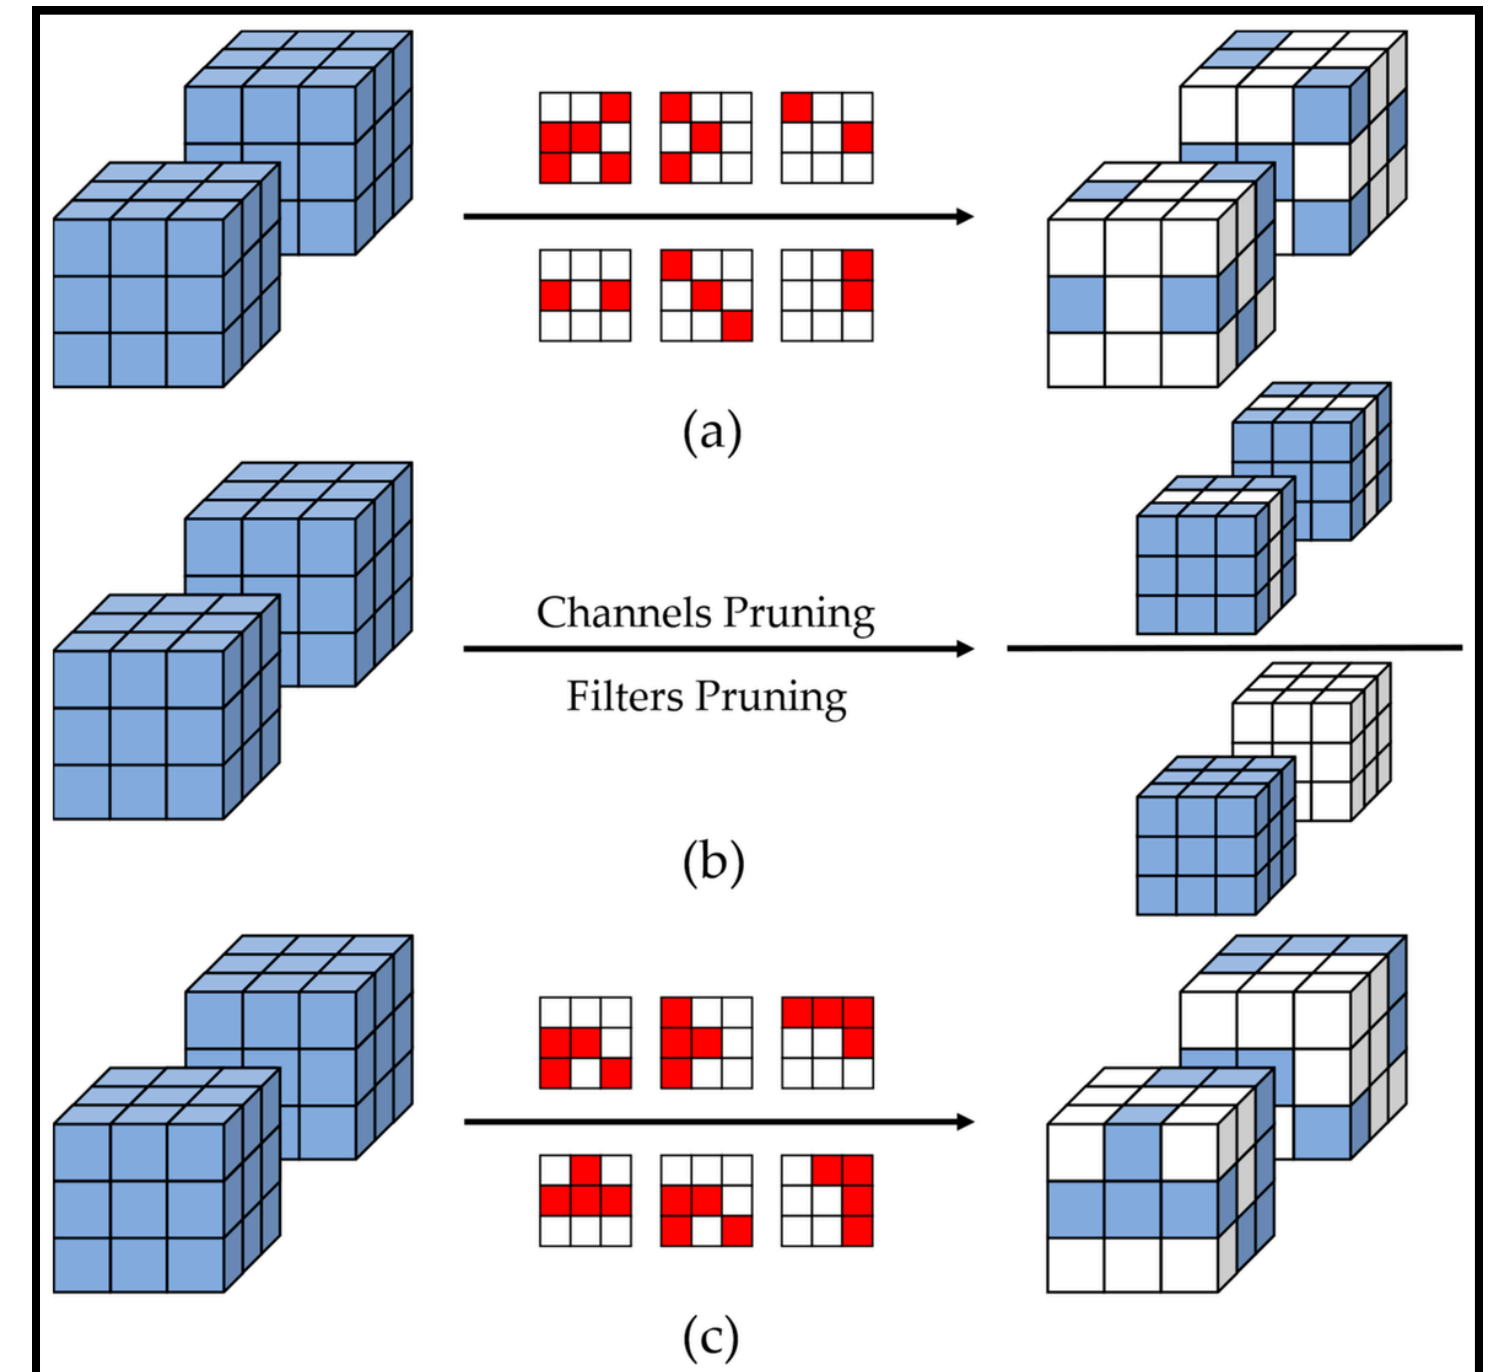
Why 50% Pruning?

Pruning Rate	Model Size	Compression Ratio	Accuracy
0% (baseline)	11.2MB	1.0x	54%
20%	9.1MB	1.23x	53.80%
40%	3.4MB	3.29x	53.42%
50%	0.358MB	31.3x	53.21%

Scope Of Improvement

Layer Sensitive Pruning

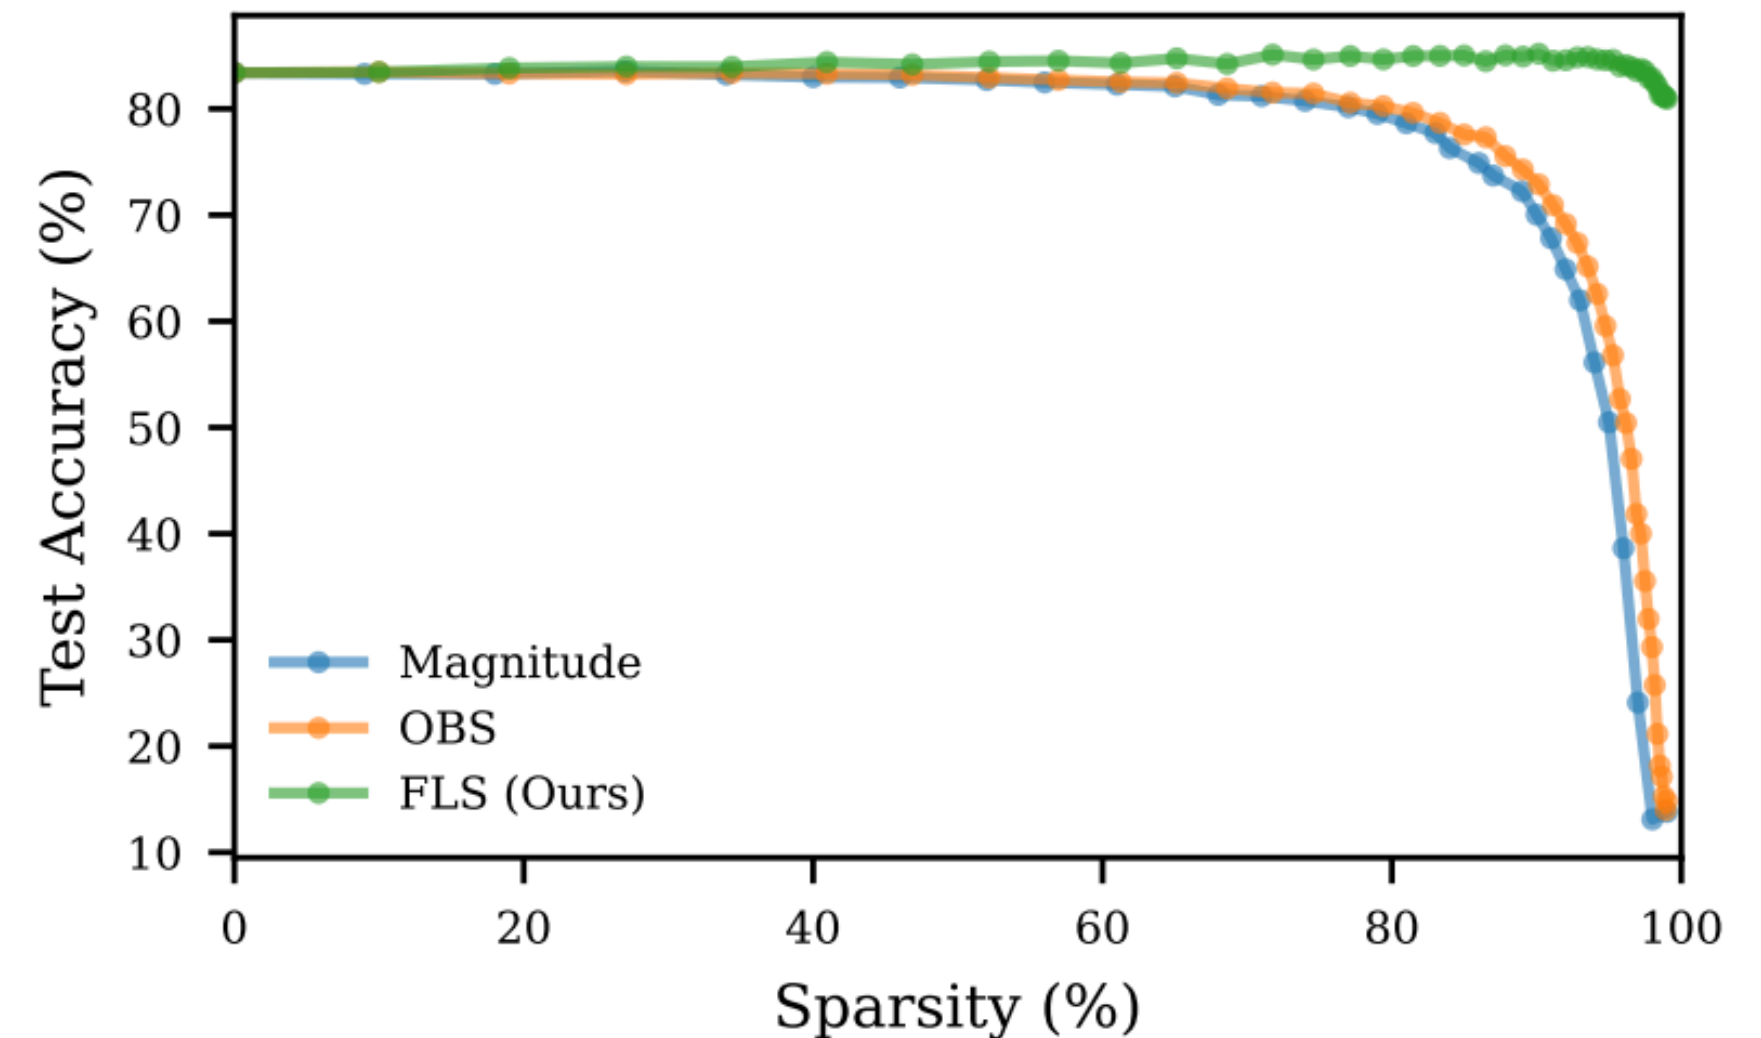
- Preserves Critical Features: Sensitive layers retain important weights, preventing accuracy drop.
- Balances Sparsity and Utility: Adapts pruning aggressiveness per layer, avoiding over- or under-pruning.
- Improves Generalization: Pruning less in high-impact layers reduces overfitting risk.
- Boosts Compression Efficiency: Achieves higher overall sparsity with minimal performance loss.



Layer Sensitive Pruning

FishLeg Pruning

- FLS builds on the FishLeg optimizer to directly learn the inverse Fisher Information Matrix (FIM) using a meta-learned parametric form $Q(\lambda)$.
- FLS uses second-order information (curvature of the loss surface) to assess true impact of removing each weight — leading to much more informed pruning.
- OBS-like second-order methods are accurate but slow — they require recomputing the Fisher matrix or Hessian repeatedly.



FishLeg Pruning

- FLS avoids this by learning a compact, structured representation of the inverse Fisher — so it saves time and memory.
- Supports unstructured and semi-structured (2:4) pruning.
- Utilizes tensor factorization (Kronecker/block-diagonal) for memory efficiency.
- Uses smart initialization ($\alpha \approx \gamma^{-1}$) and preconditioning for faster convergence.

